

NEUROGENESIS: The HELIX Based Genesis Engine of Self-Actualizing Intelligence

A Prototype for Autonomous Streaming Adaptation with Drift-Aware Replay

Saif Mansuri
Ai Engineer

[Twitter](#)

[LinkedIn](#)

[Github](#)

saifakhtar43528@gmail.com

Thakur College of Engineering and Techonology

Abstract:

Traditional AI models are trained on past data and often struggle when the data changes over time. We introduce HELIX (Hierarchical Evolutionary Learning & Intelligent eXaptation), a framework designed to continuously monitor its performance, detect changes in incoming data, and automatically choose how to adapt—through retraining, incremental learning, or ensemble expansion.

Inspired by DNA, HELIX acts as a synthetic genome by combining drift detection with a memory of past knowledge to reduce forgetting. **We test HELIX on the Allen Institute’s Developing Mouse Visual Cortex single-cell dataset, using different developmental stages as sequential data environments.** Our goal is to show that AI can adapt to changing data while preserving previously learned knowledge, with minimal human intervention.

I. INTRODUCTION: From Static Models to Self-Evolving Genomes

The deep learning revolution gave us powerful models, but they remain snapshots of a past world. When the data distribution shifts—as it inevitably does in biology, finance, or social media—these models decay. The conventional remedy is to retrain from scratch, a costly, manual, and slow process.


This article introduces **HELIX**, a framework that redefines the paradigm. Much like DNA uses a genetic code to adapt, repair, and evolve across generations, HELIX provides a synthetic genome for digital systems. It is not a single algorithm, but an architecture that embodies three core capabilities:

- **Self-awareness** – continuous monitoring of the model’s performance and the incoming data stream.
- **Self-adaptation** – upon detecting a significant change, it autonomously chooses the most appropriate update strategy (retrain, partial fit, or ensemble).
- **Self-preservation** – a memory buffer with importance-weighted replay prevents catastrophic forgetting, ensuring that past knowledge is retained.

HELIX is not a futuristic AGI—it is a working prototype, implemented in Python, that you can run today. We demonstrate it on the Allen Institute’s developmental brain dataset, where gene expression profiles change across five distinct stages (embryonic to adult). This natural sequence provides a perfect testbed for streaming adaptation.

2. Background – The Building Blocks of Self-Evolving Systems

HELIX: From Static Models to Self-Evolving Genomes



2.1 Streaming Machine Learning –

Online algorithms process data incrementally, updating models with each new instance. They are the foundation of any self-adaptive system.

2.2 Concept Drift –

When the statistical properties of the target variable change, the model must adapt. Detection methods like ADWIN (Adaptive Windowing) and DDM (Drift Detection Method) are the “sensors” of the system.

2.3 Catastrophic Forgetting –

A well-documented phenomenon where new learning overwrites old knowledge. Replay-based methods, such as reservoir sampling, are the “memory” that preserves the past.

2.4 Meta-Learning & Strategy Selection –

The final piece: the system must not only detect drift but also decide how to adapt. This is the first step towards autonomous code-rewriting—the system chooses its own update policy.

3. HELIX Architecture – The Synthetic Genome

HELIX is a pipeline that processes data in mini-batches. Its core loop consists of:

Prediction – current model predicts on new batch.

Drift Detection – ADWIN and a Kolmogorov–Smirnov test on feature distributions.


Strategy Selection – if drift is flagged, the system chooses between:

- Full retraining with replay (for large drifts).
- Partial fit (for mild drifts, if the base model supports it).
- Ensemble expansion (add a new weak learner).
- Model Update – apply the chosen strategy.

Memory Update – add samples to the replay buffer with importance weights (higher for uncertain predictions).

Evaluation – record accuracy, forgetting, and adaptation cost.

This loop runs autonomously, requiring no human intervention once initialised.



4. The Allen Institute Dataset – A Natural Streaming Environment

We use the **Developing Mouse – Visual Cortex (Dev-Mouse-Vis-Cortex-10X)** dataset from the Allen Brain Cell Atlas. It contains 568,654 single-cell transcriptomes from 53 donors, spanning ages from E11.5 to P56. The cells are annotated with 15 cell-type classes.


We preprocess:

- Filter low-quality cells.
- Log-normalise and select 2000 highly variable genes.
- Reduce to 50 principal components.


We create **five sequential environments** by grouping ages:

1. E11.5–E14.5 (early embryonic)
2. E15.5–E18.5 (late embryonic)
3. P0–P10 (early postnatal)
4. P11–P28 (late postnatal)
5. P56 (adult)

Each stage is presented as a stream of mini-batches (2000 cells each), in chronological order. This simulates a real-world scenario where data arrives over time.



HELIX is a 'synthetic genome' for machine-learning systems



5. Implementation – Step-by-Step Code

The project is structured as follows:

```
helix_self_evolving/  
├── data/  
├── src/  
│   ├── data_loader.py  
│   ├── drift_detection.py  
│   ├── memory.py  
│   ├── models.py  
│   ├── experiment.py  
│   └── utils.py  
├── notebooks/  
├── requirements.txt  
└── run_experiment.py
```

Click On the Filenames for Accessing the file:
OR

Reach at:

<https://github.com/saif-pixel1/HELIX/tree/main>

5.1 Data Preprocessing – [data_loader.py](#).

5.2 Drift Detection – [drift_detection.py](#).

5.3 Replay Memory – [memory.py](#).

5.4 Model Wrappers – [models.py](#).

5.5 Run Experiment – [run_experiment.py](#).


6. Evaluation, Metrics & Results

We measure:

Accuracy on each stage.

Forgetting = average drop in accuracy on earlier stages after training on all stages.

Adaptation count – number of drift-triggered updates.



7. Performance Summary

HELIX achieved the highest average accuracy (0.80) and the lowest forgetting (0.04) among the three approaches. Its performance remained relatively stable across developmental stages, while the Static model showed a larger decline as the data distribution changed. HELIX triggered 3 adaptations across 5 stages, indicating that it adapted selectively rather than updating unnecessarily.


8. Path to Full Autonomy – Towards Self-Rewriting Code

The current HELIX prototype chooses among fixed update strategies. However, the next evolutionary step is to allow the system to rewrite its own code – for example, by:

- **Genetic Programming** – evolving new feature extractors or model architectures.
- **Online Hyperparameter Tuning** – using Bayesian optimisation to adjust model parameters on the fly.
- **Neural Architecture Search** – dynamically expanding or pruning network layers.

These techniques are already being researched; integrating them into HELIX would transform it from a self-adaptive framework into a truly self-evolving organism. This is the direction we are currently exploring.

9. Till now Anatomy is:



9. Conclusion

HELIX embodies the vision of the notice: a synthetic genome that endows machines with the ability to sense change, adapt intelligently, and preserve knowledge – all without human intervention. Our prototype, tested on the Allen developmental brain dataset, demonstrates clear advantages over static and naive approaches. With the open-source code provided, you can replicate and extend this work, taking the first step towards building AI that does not merely process data, but actively evolves its own intelligence.

References

- Yao, Z., van Velthoven, C.T.J., Kunst, M., et al. "A high-resolution transcriptomic and spatial atlas of cell types in the whole mouse brain." *Nature*, vol. 624, no. 7991, 2023, pp. 317–332. doi:10.1038/s41586-023-06812-z.
- Allen Institute for Brain Science. "Developing Mouse - Visual Cortex single cell transcriptomes." Allen Brain Cell Atlas. Available: https://alleninstitute.github.io/abc_atlas_access/
- Allen Institute for Brain Science. "Mouse Whole Cell Tissue Processing for 10x Genomics Platform V.9." *Protocols.io*, 2022. doi:10.17504/protocols.io.q26g7b52klwz/v9.
- Read, J., and Žliobaitė, I. "Supervised Learning from Data Streams: An Overview and Update." Research Portal IP Paris, 2025.
- Garcia, C.M., Abilio, R., Koerich, A.L., Britto, A.S., and Barddal, J.P. "Concept Drift Adaptation in Text Stream Mining Settings: A Systematic Review." *ACM Transactions on Intelligent Systems and Technology*, vol. 16, no. 2, 2025, pp. 1–67. doi:10.1145/3704922.
- Mussa, A.A., Alfoudi, A.S., and Asaedi, A.H. "A Survey on Incremental Learning Techniques for Streaming Data." *Journal of Qadisiyah Computer Science and Mathematics*, 2026.
- Žliobaitė, I. "Learning under Concept Drift: An Overview." arXiv:1010.4784, 2010.
- Drift Detection Methods
- Gama, J., Medas, P., Castillo, G., and Rodrigues, P. "Learning with Drift Detection." In *Advances in Artificial Intelligence – SBIA 2004*, Springer, 2004, pp. 286–295. (DDM - Drift Detection Method)
- Bifet, A., and Gavalda, R. "Learning from Time-Changing Data with Adaptive Windowing." In *Proceedings of the 2007 SIAM International Conference on Data Mining*, 2007, pp. 443–448. (ADWIN)
- Baena-García, M., del Campo-Ávila, J., Fidalgo, R., Bifet, A., Gavalda, R., and Morales-Bueno, R. "Early Drift Detection Method." In *Fourth International Workshop on Knowledge Discovery from Data Streams*, 2006. (EDDM)
- Barros, R.S.M., Cabral, D.R.L., Gonçalves, P.M., and Santos, S.G.T.C. "RDDM: Reactive Drift Detection Method." *Expert Systems with Applications*, vol. 90, 2017, pp. 344–355. doi:10.1016/j.eswa.2017.08.023. (RDDM)
- Assis, D.N. "ADWIN-U: Adaptive Windowing for Unsupervised Drift Detection on Data Streams." *Data Mining and Knowledge Discovery*, 2025. doi:10.1007/s10115-025-02523-1.
- "Addressing Catastrophic Forgetting in Class-Incremental Learning—A Survey." *Artificial Intelligence Review*, vol. 59, article 180, 2026. doi:10.1007/s10462-026-11575-w.